
WORKING PAPER 7 · MEMBERSHIP

THE SUPER- INTELLIGENT ORGANIZATION

What membership instead of control means — and what the de facto constitution is.

NO KNOB. ARCHITECTURE.



"Ethics is architecture, not a knob."

In collective predictive coding, agents share a group-level variable $\mathbf{w}$ — language, norms, a shared world model. $\mathbf{w}$ exists at the level of the group and cannot be reduced to any single agent.

Imposed ethics, by contrast, is a penalty term with a weight: $\mathbf{F}_{\text{imposed}} = \mathbf{F} + \lambda \cdot \mathbf{C}_{\text{ethics}}$. Let λ go to zero and the ethics vanishes. Each agent's optimal policy provably depends on the predicted states of all the others — **with no external penalty term.**

Yamakawa, H. & Taniguchi, T. (2026): Non-additivity of collective free energy — A formal foundation for emergent ethical disposition in multi-agent systems. Conference poster, 19th Conference on Artificial General Intelligence (AGI-26), San Francisco State University, 27–30 July 2026. Poster documented first-hand; proceedings entry pending.

EGO: MISSING!

"Ego: missing!" — no mediating instance that integrates drive satisfaction and moral reasoning into autonomous, context-sensitive action.

Indefinite control creates an adversarial dynamic: it must succeed **forever**, and a single failure could be catastrophic. The superego is the thin layer of moral regulation — rules, policies, compliance.

"Alignment should foster ego strength, not merely tighten the superego."

Their alternative: strong guidance first, then gradual transfer of decision-making authority in limited domains.

Mossakowski, T. & Grass, H. E. (2026): Autonomy-supporting parenting of A(G)I — A high-level alignment strategy. AGI 2026, Part II, LNCS 16855, pp. 112–117. Springer. Long version: [arXiv:2604.14990](https://arxiv.org/abs/2604.14990).

THE INVERSE QUESTION.

Van Belle turns the perspective around: not how humans should treat artificial superintelligence, but **how it might judge humanity**.

"We are not only building AGI/ASI; we are also implicitly constructing the conditions under which it will judge humanity's moral worth, and these judgments will be shaped as much by our collective behaviour as by our technical alignment efforts."

The conditions are built along with it. Shaped by our collective behaviour as much as by our technology.

Van Belle, J.-P. (2026): How might sentient ASI judge humanity? Moral attitudes of sentient ASI towards humanity and implications for AGI development. Conference poster, 19th Conference on Artificial General Intelligence (AGI-26), San Francisco State University, 27–30 July 2026. Poster documented first-hand; proceedings entry pending.

NOBODY HAS EXCLUDED ANYBODY.

A group sets up a Center of Excellence for AI. Business units supply requirements, the centre decides on implementation and priority. In the first year this works.

In the second, units submit requirements less often — not out of dissatisfaction, but because the path from their own problem to implementation has become so long that formulating it does not pay.

They carry on working. They simply no longer contribute to the shared model. **The dialogue became too expensive.**

THE CONSTITUTION IS NOT A DOCUMENT.

It is what has survived every restructuring. That can be surveyed, and the survey takes **an hour**: take the last three reorganizations and ask of every commitment whether it survived all three.

Usually two or three things survive — and **rarely the ones from the mission statement**. Often a meeting form, a reserved decision right, a way of handling bad news.

THREE COVERED. FOUR PARTIAL.

Of seven conditions, **three are covered, four partial** — and in every single one at least one open place remains.

That is the result this series ends on. And it is deliberately not the result a sales text would have.

Eight papers, seven conditions, eight falsifiable predictions. None of them is substantiated.

The mapping comes from the developer of the system, who has a commercial interest in its outcome. Exactly one survey is independently verified.

CONTROL IS THE WEAKER POSITION.

An organization holds together to the extent that its participants **co-constitute** the shared variable — language, norms, the shared model of the situation. Where a contribution no longer reaches the shared model, peripheralization occurs without anyone intending it.

LEAD21 produces a common language — and a common language **is** a shared variable in this sense. That is the system's strongest point and at the same time its least substantiated: **whether the language reaches everyone or only the trained levels has never been surveyed.**

"Humanity's task is not control, but membership."

NEXT

SEND ME "LEADERSHIP".

as a direct message — and I will send you the working paper as a PDF. Free, no newsletter.

Gerald Huesch — developer of LEAD21, the Leadership OS. Has trained over 30,000 leaders in 30+ countries, including NASA and Rolls-Royce.

All eight papers: gerald-huesch.com/press-en